
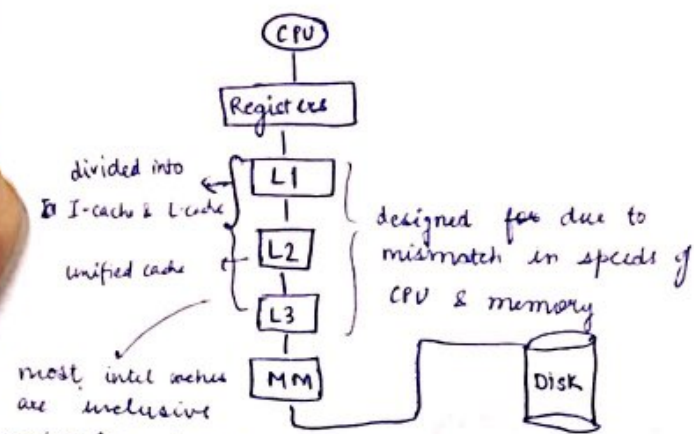





THE MEMORY HIERARCHY :






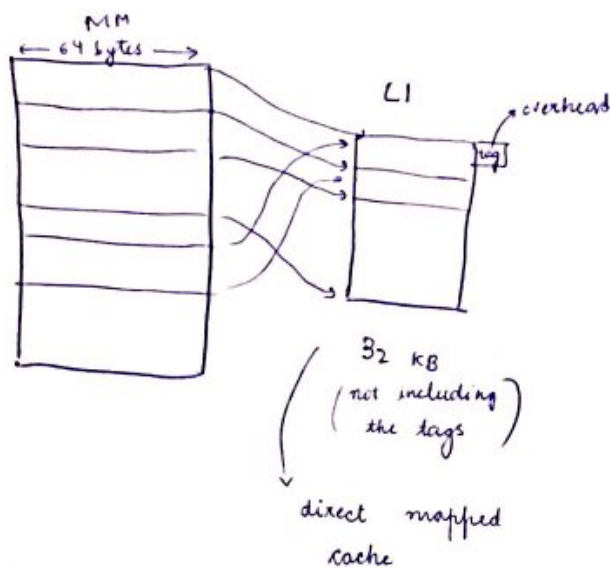
Name	Regs	cache	MM	Disk
Size	< 1KB	64 K/256 K/2.4 MB	4-64 GB	> 1TB
Tech	CMOS	SRAM	DRAM	Magnetic
Access time	15 - 3 nsec	5/3-10/10-20 ns	30-200 ns	5 ms
manage by	compiler	Hardware	OS	OS
Backed by	Cache	MM	Disk	Tape

b = block or line (64 bytes on most Intel M/c)

spatial locality → accessed something then high chance of accessing in the vicinity
temporal locality → " " in near past " " in future



→ L1 is virtually addressed, L2 & L3 are physically addressed



- Assume CPI = 1 w/o mem. stalls
- 50% of all instr. are Ls
- miss penalty = 25 cycles
- miss rate = 2%

So CPI after factoring mem. stalls =

Avg # of stalls per inst

$$= 1.5 \times 0.02 \times 25 = 0.75$$

So CPI = 1.75

$$\frac{1}{2} \times \frac{2}{100} \times 25$$

$$\frac{1}{2} \times 1 + \frac{1}{2} \times \frac{2}{100} \times 25$$

$$\frac{1}{2}$$


$$\frac{2}{100} \rightarrow \text{mem}$$

1.5 = average mem acc / Inst.

$$\frac{1}{2} \times 1$$

Design process:

- 1) C size (cache)
- 2) Line Size
- 3) Associativity (every line can map to a set & the cardinality of the set is associativity)
- 4) # levels
- 5) Replacement policy
- 6) Write update policy
- 7) Write Alloc vs Not alloc
- 8) split vs unified



$c \rightarrow$ cache size
 $b \rightarrow$ block/line size
 $a \rightarrow$ associativity

if 512 GB
 \downarrow
 $39 \rightarrow 40$

$64 \text{ KB} \quad 1 \rightarrow 4 \text{ KB}$
 $\quad \quad \quad a=8$

$$| \text{tag} | = 40 - 6 - \frac{2^{12}}{2^3 \cdot 2^6} = 27$$

$\log_2 512 \times 10^9$
 27
 $9 + 30$
 3×3

Miss Rate

Avg mem access time

Exce. time.

$\frac{IC}{CPI} \times \text{Cycle time} \rightarrow$ associativity will affect this
 since you will have to parallel search in the set.

$$\Rightarrow \text{Avg mem access time} = \text{Hit time} + \underbrace{MR}_{(\text{Miss rate})} \times \text{Miss Penalty}$$

Higher assoc \rightarrow ~~High~~ \downarrow MR (but at $a=8$ diminishing increase)

1-2 thumb rule \rightarrow if cache size is 32K then decrease in MR if we inc the cache size to 64K keeping $a=1$ or equivalently inc the associativity to 2 keeping cache size 32K.

\rightarrow Hit time = 1 cycle, Miss Penalty = 25 cycles
 Assume 64 KB cache

Assoc	Clock cycle time	MR	Mem Access Time
1	1	0.37	$1 + 0.37 \times 25 = 1.925$
2	1.36	0.31	$1.36 + 0.31 \times 25 = 2.135$
4	1.44	0.30	$1.44 + 0.30 \times 25 = 2.19$
8	1.52	0.29	$1.52 + 0.29 \times 25 = 2.245$


⇒ eg:- toy cache = 64 bytes
line size = 8 bytes

What is the miss rate after executing 100 iterations of the loop & each iteration has 20 instructions?

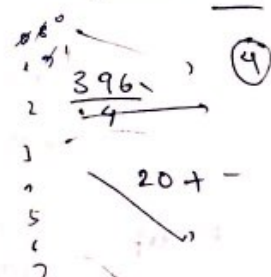

$a=1$, LRU replacement policy



20 instr



$$406 = 4 \times 99 + 10$$



In 2-way ($a=2$) →

$$6 \times 99 + 10$$

MM update policy:

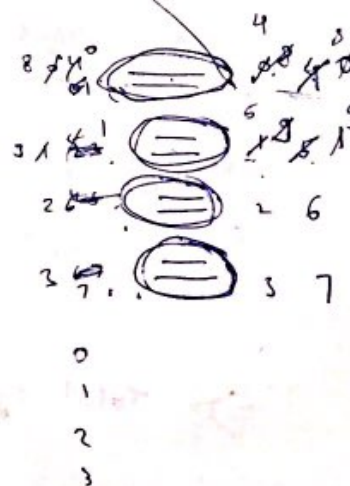
Write Through



Write Back



WRITE BUFFER



Strategy on a Write Miss

• Write Allocate

Block is allocated on a write miss

• Write Non Allocate

Block is not allocated

		W Alloc	No Alloc
W	M[100]	M	M
W	M[100]	H	M
R	M[200]	M	M
W	M[200]	H	H
R	M[100]	H	M
W	M[100]	H	H

Miss on a read
always bring it.

effectiveness of separate i-cache & d-cache

16 KB I-cache, $MR = \frac{3.82}{1000 \text{ instr.}}$

16 KB D-cache, $MR = \frac{40.9}{1000 \text{ instr.}}$

vs
32 KB unified, $MR = \frac{43.3}{1000 \text{ instr.}}$

Assume 36% instr are L/S. Instr. Hit takes 1 cycle
Data " " 1 cycle with separate D-cache
2 cycles with unified.

MP = 200 cycles

What is the AMAT?

$\Rightarrow MR_{sep} = \frac{44.72}{1360} = 0.0329$

$MR_u = \frac{43.3}{1360} = 0.0318$

AMAT

sep $\rightarrow \frac{100}{136} \times \left(1 + \frac{3.82}{1000} \times 200\right) + \frac{36}{136} \times \left(1 + \frac{40.9}{1000} \times \frac{200}{0.36}\right) = 7.57 \text{ cycles}$

\Rightarrow Total cycles = $\frac{1000}{136} + \frac{3.82}{1000} \times 200$

unified $\rightarrow 1 + 2 \times 0.36 + \frac{43.3}{1360} \times \frac{200}{0.36} = 8.08 \text{ cycles}$

$1 \times \left(2 + \frac{43.3 \times 200}{1000}\right) \times \frac{1}{100}$

11

64

(1.72) 0 (1)

(2)

3.82

$1.36 + \frac{86.64}{10} = 10.024$
 $+ \frac{43.3}{1360} \times \frac{200}{0.36}$




→

Loads	MR	MP	Access time
16	2.54	80 + 2	2.673
32	1.35	+ 4	2.139
64	1.06	+ 8	1.933
128	1.02	+ 16	1.98
256	1.15	+ 32	2.288

Projects

→ Multiprocessor cache coherence:

→



Matrix Multiplication

WB (Write Back)

WT (Write Through)

Domain Specific Archi.

→ Patterson paper (on ML)
→ Jonppi (True performance analysis)

Buffer overflow



Microarchitectural attacks

RSA , AES

↓

10 rounds - 16 table lookups & 16 XORs

4K tables (4 tables of size 1K)

AMD opteron case study:

TLB entry is 8 bytes

L1 hit time = 1 cycle

L2 hit time = 7 cycle

L1 TLB size \rightarrow 40 PTEs

L2 " " \rightarrow 512 "

\rightarrow Block repl \otimes LRU (approx)

L1 \rightarrow fully assoc

L2 \rightarrow 4 way st.

L1-split I-cache & D-cache 64 KB

2-way SA, Hit time = 3 cycles

LRU;

L2 - unified 512 KB

16-way SA

Hit time = 9 cycles

L3 - unified, 2 MB

32-way SA

Hit time = 38 cycles

I/O - related terms


• mem mapped I/O vs I/O instr.

vs

• control register / status reg.

• polling vs interrupt driven I/O

\rightarrow DMA (Direct mem Access)




when this guy is ready with disk info,
it will place it in main memory

Reliability, Availability, MTTF, MTBF, MTTR

$R(t) = \text{prob} \{ S \text{ is functioning in } [0, t] \}$
 $N_o(t)$ of N components are functional at t
 $N_f(t)$ of N components have failed in $[0, t]$

$$R(t) = \frac{N_o}{N}$$

$$\frac{dR}{dt} = -\frac{1}{N} \frac{dN_f}{dt}$$



$\rightarrow Z(t)$ (failure rate or hazard rate) $= \frac{1}{N_o} \frac{dN_f}{dt} = \left(-\frac{\frac{dR}{dt}}{R} \right)$

$\text{If } Z(t) = \lambda \text{ (on start)}$


$\text{then } R(t) = e^{-\lambda t}$

\rightarrow Let x be a r.v. representing lifetime of a component & let F be cdf of x

$\text{Then } R(t) = \text{Pr}[x > t] = 1 - F(t) = \int_t^{\infty} f(x) dx = e^{-\lambda t}$

MTTF:

$$\text{MTTF} = \int_0^{\infty} x f(x) dx = \int_0^{\infty} x \lambda e^{-\lambda x} dx = \frac{1}{\lambda}$$



$$\text{Availability (uptime)} = \frac{\text{MTTF}}{\text{MTTF} + \text{MTTR}}$$

AFR (Annual failure rate):

- % of devices expected to fail in a year for a given MTTF.
- 50K servers, each with 2 disks

$$\begin{aligned} \text{MTTF} &= 1000000 \text{ hours} \\ &= 14 \text{ years} \end{aligned}$$

$\text{AFR. Annually how many disks will fail} = \frac{8760}{10^6} \times 10^5 = 876 \text{ disks}$

2 disks failing everyday